Interpretability

Natural language processing

Student: Ali Dankoob

40407024

Papers list

File Edit View Insert Format Data Tools Extensions Help

Share

Menus 100% Comment only

A16:E16 Model Interpretability

	A	B	C	D	E	F	G	H
1	Paper_ID	کلان‌موضوع	عنوان مقاله	لینک	نوع مسئله	نام دانشجو	شماره دانشجویی	ایمیل
2	1	Rumor/LLM	Enhancing large language model capabilities for rumor detection with Knowledge Powered Prompting	https://www.sciencedir	Rumor Detection (Classification)			
3	2	Rumor/LLM	HRDE: Retrieval-Augmented Large Language Models for Chinese Health Rumor Detection and Explainability	https://arxiv.org/abs/24	Rumor Detection (Retrieval-Augmented)			
4	3	Stance/Farsi	FarExStance: Explainable Stance Detection for Farsi	https://aclanthology.or	Stance Detection			
5	4	Emotion/Farsi	Emotion Alignment: Discovering the Gap Between Social Media and Real-World Sentiments in Persian Tweets and Images	https://arxiv.org/html/2	Emotion/Sentiment Analysis			
6	5	Topics/Emergen	Evaluation of Unsupervised Static Topic Models' Emergence Detection Ability	https://doi.org/10.7717	Topic Modeling / Emergence Detection			
7	6	Concept Extracti	ConExion: Concept Extraction with Large Language Models	https://arxiv.org/pdf/25	Concept Extraction			
8	7	Concept Extracti	Data-Efficient Concept Extraction from Pre-trained Language Models	https://aclanthology.or	Concept Extraction			
9	8	Concept Extracti	Causality-aware Concept Extraction based on Knowledge-guided ...	https://aclanthology.or	Concept Extraction			
10	9	Concept Extracti	A Weakly Supervised Approach Using Contextualized Word Embeddings	https://aclanthology.or	Concept Extraction (Weakly Supervised)			
11	10	Event Extraction	An Ensemble Event Extraction Method on News	https://link.springer.cor	Event Extraction			
12	11	Opinion Concept	Unified Opinion Concepts Ontology and Extraction Task	https://aclanthology.or	Opinion Concept Extraction			
13	12	NER/News	Named Entity Recognition (NER) for News Articles	https://doi.org/10.3421	NER for News			
14	13	Topic Modeling	Semantic-Driven Topic Modeling Using Transformer-Based ...	https://arxiv.org/pdf/24	Transformer-based Topic Modeling			
15	14	Topic Modeling	Probabilistic Topic Modelling with Transformer Representations	https://arxiv.org/pdf/24	Probabilistic Topic Modeling			
16	15	Interpretability	Linguistic Interpretability of Transformer-based Language Models	https://arxiv.org/pdf/2	Model Interpretability			
17	16	Topics/Transforn	Disentangling Transformer Language Models as Superposed Topic ...	https://aclanthology.or	Topic Analysis in Transformers			
18	17	AI-Generated Te	A Transformer-based Framework for Detecting AI-Generated ...	https://aclanthology.or	AI-generated Text Detection			
19	18	Causal LM	A Meta-Learning Perspective on Transformers for Causal Language ...	https://aclanthology.or	Causal Language Modeling			
20	19	ABSA	Aspect-based Sentiment Analysis via Synthetic Image Generation	https://aclanthology.or	Aspect-Based Sentiment Analysis			
21	20	ABSA	Exploring Aspect-Based Sentiment Analysis Methodologies	https://aclanthology.or	ABSA (Methodology Review)			
22	21	ABSA	Instruction Learning for Aspect Based Sentiment Analysis	https://aclanthology.or	ABSA with Instruction Learning			
23	22	ABSA	Improving Aspect-based Sentiment Analysis with Diffusion Models	https://aclanthology.or	ABSA with Diffusion Models			
24	23	ABSA	Instruct-DeBERTa: A Hybrid Approach for Aspect-based Sentiment ...	https://arxiv.org/pdf/24	Hybrid ABSA (Instruct + DeBERTa)			
25	24	Aspect Extractio	An Aspect Extraction Framework Using Different Embedding Types	https://arxiv.org/pdf/25	Aspect Extraction			
26	25	Scenario/Decisic	Optimal Decision Making Through Scenario Simulations Using Large Language Models	https://arxiv.org/abs/24	Scenario Simulation / Decision Support			
27	26	Agent Framewor	STRIDE: A Tool-Assisted LLM Agent Framework for Strategic and Interactive Decision-Making	https://www.arxiv.org/	LLM Agent Framework for Decision-Making			
28	27	Wargame/LLM	Simulating Open-Ended Wargames with Large Language Models (Snow Globe)	https://arxiv.org/abs/24	LLM-based Wargaming			
29	28	Agent/Policy	Exploring the Opportunities and Challenges of Using Large Language Models to Represent Institutional Agency in Land System Modelling	https://esd.copernicus.	LLM Agents as Policy-makers (ABM)			
30	29	Agent-Based Mo	Large Language Models Empowered Agent-Based Modeling and Simulation	https://www.nature.cor	Survey on LLM-powered ABM			
31	30	Human Simulati	Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies	https://arxiv.org/pdf/22	LLM-based Human Simulation			
32	31	Generative Agen	Generative Agents: Interactive Simulacra of Human Behavior	https://arxiv.org/abs/23	LLM-based Generative Agents / Social Simulation			
33	32	Strategic ToM	ToMPO: LLM-Driven Strategic Decision-Making Through Self-Play with Theory of Mind	https://openreview.net	Strategic Self-Play with Theory of Mind			

Ali Dankoob
10:47 AM Dec 24
dankoob94@gmail.com

99+ Sheet1

Sum: 15

Introduction

- Linguistic Interpretability of Transformer-based Language Models
- Graduate Course: Natural Language Processing

- Transformer-based language models such as BERT, GPT, and T5 have become the dominant paradigm in modern NLP

- Despite their strong empirical performance, their internal decision-making mechanisms remain largely opaque

- This presentation focuses on understanding how linguistic knowledge is represented and utilized inside these models

- The project is conducted in the context of interpretability and explainable artificial intelligence (XAI)

Background and Motivation

- Natural Language Processing has evolved from rule–based and statistical models to deep neural architectures
- Transformers introduced self–attention, enabling effective modeling of long–range dependencies
- This architectural shift significantly improved performance across NLP tasks
- However, it also increased model complexity and reduced transparency
- As models grow larger, understanding their internal behavior becomes increasingly critical

Problem Definition

- Linguistic interpretability studies how linguistic knowledge is encoded within transformer–based models
- It focuses on internal representations such as embeddings, layers, attention heads, and neurons
- The goal is to map abstract neural representations to human–understandable linguistic concepts
- Linguistic properties include syntax, semantics, morphology, and discourse–level information
- The problem addresses the gap between high performance and limited explainability

Scope of the Problem

- The scope goes beyond evaluating model outputs or accuracy metrics
- It investigates how language understanding emerges inside the model
- A key question is whether models truly learn linguistic structure or rely on statistical shortcuts
- Another concern is distinguishing stored linguistic knowledge from knowledge actually used in predictions
- This scope is relevant for both theoretical research and real–world deployment

Model Interpretability

- Model interpretability refers to the ability to explain a model's behavior in human terms
- It aims to answer why a specific output or prediction was generated
- Interpretability supports trust, transparency, and accountability
- In deep neural networks, interpretability is challenging due to non–linearity and high dimensionality
- Nevertheless, it is essential for debugging, validation, and responsible AI

Importance of Linguistic Interpretability

- Linguistic interpretability is especially important because language models learn linguistic patterns implicitly, without explicit supervision or formal grammatical rules. Without interpretability, it is difficult to determine whether a model's success is based on meaningful linguistic generalization or on exploiting spurious correlations in data. By revealing how linguistic structures are represented, interpretability helps identify biases, improve generalization, and ensure that models behave reliably in unseen or sensitive scenarios.

Real-world Relevance

- In real–world applications, the need for interpretability is not merely theoretical. Content moderation systems must justify why certain texts are flagged or removed. Legal and medical NLP systems require transparent reasoning to support high–stakes decisions. Educational technologies also benefit from interpretable feedback that can be understood by learners. In all these domains, linguistic interpretability increases user trust and enables accountability.

Position within NLP

- Linguistic interpretability occupies a central position within modern NLP research, bridging deep learning, linguistic theory, and explainable AI. Rather than focusing solely on performance improvements, this field aims to provide scientific insight into how language models function internally. By integrating linguistic concepts with neural analysis, interpretability research contributes to transforming NLP from a purely engineering–driven discipline into a more explanatory and theory–informed field.

Relation to Core NLP Tasks

- Most core NLP tasks rely on linguistic knowledge at different levels. Tasks such as syntactic parsing and part–of–speech tagging depend on structural understanding, while semantic tasks such as named entity recognition and semantic role labeling require meaning representation. Higher–level applications, including natural language inference, summarization, and dialogue systems, combine multiple linguistic competencies. Linguistic interpretability enables systematic analysis of how transformer models handle these tasks internally.

Research Landscape

- Research in linguistic interpretability can be broadly categorized into several complementary streams. These include probing–based approaches that examine encoded linguistic features, attention–based analyses that study structural alignment, causal methods that test the necessity of internal components, and mechanistic approaches that aim to reverse–engineer computation circuits. Together, these streams provide a comprehensive view of how linguistic knowledge emerges in transformer models.

Probing based Methods

- Probing methods train lightweight classifiers on frozen model representations
- They evaluate whether linguistic properties can be extracted from specific layers
- Common probing tasks include POS tagging, dependency relations, and semantic roles
- Layer−wise probing reveals where different types of linguistic information emerge
- A key limitation is that probing measures correlation rather than causation

Concrete Examples of Linguistic Interpretability

- Example 1: Attention Visualization
- Sentence: "The cat sat on the mat."
- Show attention heatmap from a BERT layer — highlight how cat attends strongly to sat (subject–verb relation).

- Example 2: Probing for Part–of–Speech
- Use a simple classifier to extract POS tags from different BERT layers — show accuracy per layer (e.g., lower layers capture surface forms, higher layers capture syntax).

Attention based Analysis

- Attention analysis investigates how self−attention mechanisms distribute focus across tokens
- Researchers examine whether attention heads correspond to syntactic or semantic relations
- Some heads specialize in capturing subject–verb or modifier–head dependencies
- These patterns suggest structural linguistic knowledge
- However, attention weights do not always reflect causal importance

Causal Interpretability Methods

- Causal interpretability aims to determine whether specific internal components are necessary for a model's linguistic capabilities. This is achieved through interventions such as ablation of layers, attention heads, or neurons, as well as activation patching techniques that manipulate internal states. If modifying a component leads to significant performance degradation, it can be considered causally important. These methods offer stronger interpretability claims than purely correlational analyses.

Mechanistic Interpretability

- Mechanistic interpretability focuses on understanding exact computational pathways
- It attempts to identify circuits responsible for specific linguistic capabilities
- This approach is similar to reverse engineering neural networks
- It is often applied to smaller models or constrained tasks
- Recent work explores extending this approach to large language models

Tools & Libraries for Interpretability

- Captum (PyTorch): Model-agnostic attribution methods.
- Transformers Interpret (Hugging Face): Attention and gradient-based explanations.
- LIT (Language Interpretability Tool): Interactive exploration of model predictions.
- Ecco (for GPT-like models): Visualization of generation and attention.
- Why it matters: These tools lower the barrier to entry and support reproducible interpretability research.

Key Challenges

- One of the main challenges in linguistic interpretability is that linguistic knowledge is distributed across many parameters rather than localized in individual neurons. This makes it difficult to associate specific components with clear linguistic concepts. Additional challenges include dependence on linguistic theory assumptions, high computational costs, and the absence of standardized evaluation benchmarks for interpretability research.

Limitations of Current Approaches

- Many interpretability methods rely on statistical correlations
- Results may not generalize across models or languages
- Linguistic categories may not align cleanly with neural representations
- Human interpretation introduces subjectivity
- Scalability remains a major concern for LLMs

Future Research Directions

- Future research in linguistic interpretability is expected to focus on developing more robust and scalable causal methods. Cross–lingual and multilingual interpretability will play an increasingly important role as models are deployed globally. Integrating linguistic theory into model architectures and incorporating human–in–the–loop evaluation frameworks are also promising directions. Interpretability for large language models remains a key open challenge.

Practical Applications

- In practical settings, linguistic interpretability supports model debugging, error analysis, and bias detection. It helps practitioners assess dataset quality and understand model failures. Interpretability is also essential for ensuring robustness and for supporting responsible and ethical deployment of NLP systems in real–world environments.

Interpretability for Model Improvement

- Error Analysis & Debugging: Use attention patterns to detect why NER fails on certain entities.Bias Detection: Identify stereotypical associations via probing (gender bias in coreference resolution).Data Augmentation: Use interpretability to find underrepresented linguistic phenomena and enrich training data.Interpretability-Guided Fine-tuning: Regularize models to align attention with linguistic heuristics (syntactic trees).Impact: Turns interpretability from a diagnostic into a tool for building better models.

Social Media, Content Analysis, and Security

- Explainable hate speech and toxicity detection
- Misinformation and propaganda analysis
- Discourse and narrative structure analysis
- Detection of extremist and malicious communication
- Transparent automated moderation and monitoring systems

Conclusion

- Linguistic interpretability is a core challenge in modern NLP
- It bridges neural computation and human language understanding
- Interpretability enhances transparency, trust, and safety
- The field is essential for the future of language technologies
- It plays a critical role in responsible deployment of NLP systems

References

- 1. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017).Attention Is All You Need. Advances in Neural Information Processing Systems (NeurIPS).

  2. Belinkov, Y., & Glass, J. (2019).Analysis Methods in Neural Language Processing: A Survey. Computational Linguistics.

  3. Tenney, I., Das, D., & Pavlick, E. (2019).BERT Rediscovers the Classical NLP Pipeline. Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL).

  4. Hewitt, J., & Manning, C. D. (2019).A Structural Probe for Finding Syntax in Word Representations. Proceedings of NAACL-HLT.

  5. Conneau, A., Kruszewski, G., Lample, G., Barrault, L., & Baroni, M. (2018).What You Can Cram into a Single Vector: Probing Sentence Embeddings for Linguistic Properties. Proceedings of ACL.

  6. Hewitt, J., & Liang, P. (2019).Designing and Interpreting Probes with Control Tasks. Proceedings of EMNLP-IJCNLP.

7. Clark, K., Khandelwal, U., Levy, O., & Manning, C. D. (2019).What Does BERT Look at? An Analysis of BERT's Attention. Proceedings of ACL.

8. Jain, S., & Wallace, B. C. (2019).Attention Is Not Explanation. Proceedings of NAACL-HLT.

9. Elhage, N., et al. (2021).A Mathematical Framework for Transformer Circuits. Transformer Circuits (OpenAI / Anthropic).

10. Manning, C. D., Clark, K., Hewitt, J., Khandelwal, U., & Levy, O. (2020).Emergent Linguistic Structure in Artificial Neural Networks. Proceedings of the National Academy of Sciences (PNAS).

11. Doshi-Velez, F., & Kim, B. (2017).Towards a Rigorous Science of Interpretable Machine Learning. arXiv preprint arXiv:1702.08608.

12. Danilevsky, M., Qian, K., Aharonov, R., Katsis, Y., Kawas, B., & Sen, P. (2020).A Survey of the State of Explainable AI for Natural Language Processing. Proceedings of ACL.

The End

Thank you for your attention